

CS4361/5361 Machine Learning

Fall 2019

Lab 3 - Decision and regression trees

Due Monday, October 7, 2019. Submit a paper copy by 5:00 p.m.

Email report to *olacfuentes@gmail.com*, include UTEP-ML2019 in the subject line.

Your task for this lab is to extend your implementation of decision and regression trees and to implement three algorithms for generating ensembles.

Part 1

Extend the decision tree functions as follows:

1. Write a function to remove unnecessary internal nodes from your decision tree. A node is unnecessary if both of its children are leaves and they result in the same classification. Perform a post-order traversal on your tree, replacing each unnecessary internal node by the corresponding classification. To test your implementation, make sure the accuracy before and after pruning is the same.
2. Write functions to determine the number of internal nodes in your decision and regression trees.
3. Write functions to determine the height of your decision and regression trees.
4. Write a function to determine the importance of each attribute in the dataset from the structure of the tree. A good measure of an attribute's importance is the average number of times it is used to predict the target function value of each training example.
5. (Extra credit) Optimize your prediction time by eliminating the loop in the predict function. This will require to modify the classify function as well.

Part 2

Implement the following ensembles of predictors and evaluate their performance:

1. Randomization
2. Bagging
3. Boosting

Write a report including (at least) the following items:

1. Problem description
2. Algorithms implemented
3. Experimental results. It is important that you try to obtain the best possible results by optimizing the parameters of your trees and ensembles. Show accuracies or mean squared errors and running times for each algorithm and parameter choice. You may want to use tables and or plots to illustrate this.
4. Discussion of results. What ensemble and parameter selection yields the best results? Why?
5. Conclusions
6. Appendix: Source code

Appendix

Sample run:

Original Tree:

```
        class= 1
      if x[0] <= 43.3278
        class= 1
      if x[6] <= 24.0087
        class= 1
      if x[3] <= 0.238304
        class= 1
    if x[8] <= 9.49499
      class= 1
      if x[1] <= 14.6079
        class= 1
      if x[0] <= 45.6028
        class= 1
        if x[0] <= 86.9231
          class= 0
    if x[8] <= 27.6336
      class= 1
      if x[2] <= 2.30643
        class= 1
      if x[2] <= 2.50051
        class= 0
      if x[1] <= 16.3897
        class= 1
    if x[0] <= 53.3287
      class= 0
      if x[9] <= 198.209
        class= 0
      if x[1] <= 49.4408
        class= 0
      if x[5] <= -63.1967
        class= 0
```

Number of internal nodes: 15

Height: 4

train accuracy: 0.80744

test accuracy: 0.81519

Pruned Tree:

```
        class= 1
      if x[8] <= 9.49499
        class= 1
      if x[0] <= 45.6028
        class= 1
        if x[0] <= 86.9231
          class= 0
    if x[8] <= 27.6336
      class= 1
      if x[2] <= 2.50051
        class= 0
      if x[1] <= 16.3897
        class= 1
    if x[0] <= 53.3287
      class= 0
```

Number of internal nodes: 7

Height: 4

Attribute usage per example:

Attribute: 0 used 0.7169 times per training example

Attribute: 1 used 0.1316 times per training example

Attribute: 2 used 0.2780 times per training example

Attribute: 3 used 0.0000 times per training example

Attribute: 4 used 0.0000 times per training example

Attribute: 5 used 0.0000 times per training example

Attribute: 6 used 0.0000 times per training example

Attribute: 7 used 0.0000 times per training example

Attribute: 8 used 1.6120 times per training example

Attribute: 9 used 0.0000 times per training example

Most important attribute: 8

train accuracy: 0.80744

test accuracy: 0.81519